

PROPOSAL · AGENTIC AI CAPABILITY

From Conversation to Action.

InTrust that understands, decides and completes.


InTrust already speaks with customers and merchants. This proposal extends that conversation into an **agentic capability** — one that understands intent, retrieves only permitted context, selects an approved capability, and completes the request through the workflows InTrust already runs.

- SCOPED TO THE SIGNED-IN USER
- APPROVED CAPABILITIES ONLY
- EXISTING WORKFLOWS STAY AUTHORITATIVE
- PAYMENT STAYS OUTSIDE THE AGENT

Prepared for	InTrust India	Experiences	Customer · Merchant
Initiative	Agentic AI	Status	Proposal

PROPOSED EXPERIENCE

one ask. one action.



✓

Action completed — cart updated

DONE

Executed by the existing InTrust cart workflow



Text, image, voice

Multimodal requests understood as a single intent.



Two experiences

Customer commerce and merchant operations.



One intelligence layer

Added on top of the platform already in use.



No payment access

Financial execution stays inside secure checkout.

INSIDE THIS PROPOSAL

✓

The shift. Why conversational replies become completed actions.

✓

The reasoning. Context, decision flow and capability selection.

✓

The architecture. How an agentic layer sits over existing systems.

✓

The governance. Permissions, guardrails and the action model.

✓

The experiences. Customer journeys and the merchant assistant.

✓

The rollout. A phased path from foundation to automation.

GUARDRAILS AT A GLANCE

the model proposes

✓

Server-side identity — the agent never decides who you are.

✓

Registered capabilities — no arbitrary queries, no open API access.

✓

Existing logic decides — the platform remains the authority.

5 THE AGENTIC LOOP, IN ONE LINE

understand · decide · act

1

Understand

Intent extracted from text, image or voice.

→

2

Retrieve

Only permitted context is assembled.

→

3

Decide

An approved capability is selected.

→

4

Execute

Existing InTrust logic performs the action.

→

5

Return

A confirmed result returns to the user.

“ InTrust already has the workflows. Agentic AI gives its users a faster, more natural way to reach them — without weakening a single control. ”

3

INPUT TYPES

2

EXPERIENCES

1

PLATFORM LAYER

0

PAYMENT ACCESS

EXECUTIVE VISION

the shift that matters
INTRUST · AGENTIC AI

From smarter replies to completed requests.

InTrust already captures demand at the moment of intent — in chat, on the storefront, in the merchant's daily operations. Today that conversation ends with information. The opportunity is to let it end with a **result**: an item added, an order explained, a merchant question answered from their own data, a cart prepared for a secure checkout.


Today — the assistant answers

- Answers a question and stops at information.
- Explains steps the user must then perform manually.
- Hands the user off to a sequence of screens.
- Every journey restarts from the beginning.

→


Proposed — the agent completes

- Understands intent from text, image, voice or a mix.
- Assembles only the context the signed-in user is allowed to see.
- Selects an approved capability from a fixed registry.
- Completes the request through the workflows InTrust already runs.


Fewer steps to the same outcome

Natural language replaces multi-screen navigation, so demand is acted on while it is still fresh.


The trust model stays intact

Authentication, permissions and business rules remain the authority. The agent operates inside them.


One foundation, two audiences

The same intelligence layer serves customers and merchants through different permissions.

DIRECTIONAL GOAL
Fewer taps per request
 Baseline to be agreed in discovery.

DIRECTIONAL GOAL
Higher intent completion
 Measured on existing funnel events.

DIRECTIONAL GOAL
Faster merchant answers
 From panel navigation to one question.

DIRECTIONAL GOAL
Consistent policy
 Every action passes the same checks.

SCOPE NOTE Nothing in this document should be read as delivered functionality. It describes a **proposed** capability, and the design principles it would follow if built.


User-scoped
 Every read and action is limited to the caller.


Capability-bounded
 Only reviewed actions from a fixed registry.


Server-authoritative
 The platform decides, not the model.


Observable
 Requests, selections and outcomes are traceable.


Reversible
 Any capability can be switched off independently.


Questions we would baseline with InTrust

- Which customer and merchant journeys matter most today?
- Which events already measure completion, and what are their current values?
- Which actions feel safe enough to execute without a confirmation step?
- What support and merchant feedback themes should the agent absorb first?


Commitments this proposal makes

- No new source of truth, and no duplicate commerce backend.
- No change to how customers authenticate or check out.
- No capability enabled until it has been reviewed and approved.
- No silent state changes — every outcome is visible to the user.

NOT IN SCOPE No new app to install, no data migration and no change to how InTrust customers pay — and **no decision needed today** beyond agreeing the shape of the first phase.

“ A reply informs the customer. An agent changes the state of their account — safely, verifiably, and only within what they are permitted to do. ”


Fewer
 STEPS TO ACT









Same
 TRUST MODEL


Same
 PLATFORM

WHAT AGENTIC AI MEANS HERE

not a new app — a new way in

Chatbots answer. Agents **act**.




AGENTIC ARCHITECTURE

One intelligence layer. Zero parallel systems.


an action layer, not a rebuild

The agent sits **between** the user's conversation and the services InTrust already operates. It interprets the request, assembles permitted context, chooses from a fixed set of registered capabilities, and hands execution to the platform. Every result travels back through the same path in reverse — so the conversation never becomes a second source of truth.



6 THE REQUEST PATH, EVERY TIME

same checks, same rules



+ What the layer adds

NEW

- ✓ Intent understanding across text, image and voice.
- ✓ Context assembly limited to the caller's permissions.
- ✓ Capability selection from an explicit, reviewed registry.
- ✓ Guardrails, rate limits and loop protection on every request.
- ✓ Traceability: request, selection and outcome are all recorded.

What stays exactly as it is

REUSED

- ✓ **Authentication** — existing sessions and sign-in flows.
- ✓ **Permissions** — the same ownership and access rules.
- ✓ **Business logic** — pricing, availability, eligibility, fulfilment.
- ✓ **Data** — one catalogue, one cart, one order record.
- ✓ **Audit** — the platform's own logs remain authoritative.



No duplicate backend

The agent calls services, it does not recreate them.



Observable by default

Each hop can be inspected without changing the platform.



Reversible

Capabilities can be enabled or withdrawn independently.



Integration surface — proposed

- ✓ The agent calls **existing service endpoints**, not raw tables.
- ✓ Identity and permissions are resolved **server-side** from the session.
- ✓ Read-only context is enabled first; write actions follow capability by capability.
- ✓ No schema change is required to add the layer.



Behaviour when something fails

- ✓ The user is told what could not be completed, in plain language.
- ✓ No partial state is hidden — the conversation reflects reality.
- ✓ Retries are bounded; the agent never loops indefinitely.
- ✓ Every failure is logged with the capability that was attempted.

DESIGN RULE

If a request cannot be served by a **registered capability**, the agent explains the limitation and routes the user to the normal journey. It never improvises a workaround.

“ **Interfaces do not change. The catalogue does not change. The order journey does not change. What changes is how quickly a user can reach them.** ”



1
ADDED LAYER



1
SOURCE OF TRUTH




6
REUSED SERVICES

CUSTOMER EXPERIENCE

ask once, done once

The customer stops navigating. The assistant starts completing.



Product discovery

"Show me...", "Find something like...", "Under ₹3,000".



Search & filtering

Category, price, size and availability in one sentence.



Visual discovery

An image becomes a search — style, colour or category.



Wishlist

Add, review or remove saved items conversationally.



Cart actions

Add, update quantity and confirm contents and totals.



Order lookup

Status, delivery estimate and history for the signed-in user.



Rewards & wallet

Balance and tier read from existing systems.



Checkout preparation

Cart prepared and confirmed — the secure flow still pays.



Every request is scoped to the caller

The agent can only read and act on data belonging to the authenticated customer. There is no path from one account to another.

3 WHAT CHANGES FOR THE CUSTOMER

Concept mockup · proposed experience, not delivered functionality

BEFORE

Search, tap, scroll, repeat

The customer interprets the catalogue and performs every step.



PROPOSED

One clear request

The intent is expressed in the customer's own words.



RESULT

A confirmed outcome

The cart, wishlist or order question is already resolved.



Fewer abandonment points

Intent is captured and acted on inside the conversation, instead of being handed back to the customer.



Nothing is hidden

Prices, quantities and availability are always read from the platform, shown before anything is final.



One confirmation standard

The same rules apply whether the action starts in the app or in a sentence.

“ Conversation removes the navigation, not the customer's control. They still see the outcome, the price and the confirmation before anything is final.



8

CAPABILITY
FAMILIES



1

ACCOUNT
SCOPE





0

PAYMENT
ACCESS

MULTIMODAL INTELLIGENCE

say it however you like

A photo, a sentence, a spoken word — the same intent.



Why multimodal matters here

- ✓ Customers describe products far better with a picture than with words.
- ✓ Voice removes typing friction on mobile, where most traffic lives.
- ✓ A single intent layer means no new journey to design per input type.

Boundaries that do not move

- ✓ An uploaded image is used only to interpret this request.
- ✓ Voice notes follow the same data-handling rules as text.
- ✓ No modality can bypass a permission check or a confirmation step.

“ A picture is often a better request than a sentence. Multimodal understanding lets InTrust accept whichever one the customer actually has. ”

4
INPUT MODES

1
INTENT LAYER


1
PERMISSION GATE

AGENTIC COMMERCE JOURNEY

buying, without the maze

From “I want this” to a prepared checkout.

The commerce journey is where an agent earns its place — and where the boundaries matter most. The agent takes the customer to a **review-ready cart**; the platform keeps control of the money. Every step below maps to a capability InTrust already has today.



CART PREPARED · REVIEW READY 2 items


Aero Glide — UK 9
Qty 2 · In stock

₹5,798


Sports bottle — 750 ml
Qty 1 · In stock

₹599

Cart total · rewards applicable at checkout

₹6,397

Proceed to secure checkout

Payment, OTP and receipts are handled by InTrust's existing flow — never by the agent.

 **What the agent may do**

- ✓ Find and compare products from the live catalogue.
- ✓ Check availability, size and price before suggesting anything.
- ✓ Add, update or remove items in the customer's own cart.
- ✓ Summarise the cart and hand over to checkout.

 **What stays outside it**

- ✗ No card, wallet or UPI detail is ever seen or stored by the agent.
- ✗ No payment authorisation, OTP or 3-D Secure step is triggered by the agent.
- ✗ No order can be placed without the customer's own confirmation in checkout.

→ WHO ACTS AT EACH STEP			
STEP	AGENT'S ROLE	PLATFORM'S ROLE	WHAT THE CUSTOMER SEES
Ask	Understand the request	Provide the signed-in session	Their own words, reflected back
Identify	Choose the discovery capability	Run catalogue search and filters	Relevant products with prices
Validate	Request confirmation of details	Return price, stock and eligibility	Accurate availability information
Prepare	Call the cart capability	Perform the cart change	Updated cart contents and total
Review	Summarise and hand over	Hold the authoritative cart	A cart ready for checkout
Pay	Not involved	Authenticate and capture payment	The familiar secure payment flow



Timing stays with the user

The customer decides when to move to checkout. The agent never pushes a purchase.



Nothing final without confirmation

Cart contents, quantities and totals are always shown before the customer acts.



Same cart, same rules

Whether items arrive by tapping or by talking, one cart and one set of rules apply.

“ The agent prepares. The platform pays. A conversational route to a ready cart, with money moving only through the flows InTrust already secures.


5
AGENT STEPS




1
PAYMENT GATE


0
PAYMENT ACCESS

MERCHANT AI ASSISTANT

The merchant's whole business, one question away.

your business, in your words



Order operations
Counts, status and delivery progress.

Inventory insight
Low stock, out of stock, reorder prompts.

Business metrics
Revenue, order values, repeat customers.

Business & subscription
Plan, renewal and account status.

Serviceability
Delivery coverage for the merchant's areas.

Operational answers
Explanations of processes and policies.

Real questions the assistant answers

- ✓ "What did I earn last week?"
- ✓ "Show me the orders that need dispatch."
- ✓ "Which item sells best on weekends?"

Access is scoped to one business

- ✓ Only the authenticated merchant's own data — never another seller's.
- ✓ Read-first by default — write actions require separate approval.
- ✓ Panel permissions still apply — the assistant cannot exceed them.

Ask, then act

The assistant answers from the merchant's own records before it ever suggests a change.

Panel parity

Anything it can answer, the panel can still do — the assistant is a second door, not a replacement.

Merchant trust

A seller's data is never pooled, compared or exposed to another seller, under any phrasing.

“ Merchants should not have to learn a dashboard to understand their business. They should be able to ask it a question — and get a straight answer. ”

1
BUSINESS
SCOPE

Read
FIRST BY
DEFAULT

Same
PANEL
PERMISSIONS

CONTEXT & REASONING

remember only what is yours

Reasoning is only as good as the context behind it.

Context assembled per request

- Conversation so far — what the user actually said.
- Permitted profile context, already available in the account.
- Permitted order, cart and wishlist history for that account.
- Live catalogue, stock and price at the moment of asking.




Decision reached — and explainable

- ✓ One intent, one capability, chosen from the registry.
- ✓ The reason for the selection can be described to the user.
- ✓ Every factual value comes from a platform response, not a guess.
- ✓ If context is insufficient, the agent asks instead of assuming.

Context it never uses

- ✗ Data belonging to another customer or merchant.
- ✗ Anything the signed-in user is not permitted to see in the panel or app.
- ✗ Raw database tables or internal endpoints outside the capability list.
- ✗ Assumptions about price, stock or eligibility that no service confirmed.
- ✗ Anything the model “remembers” outside the current session's approved context.

6 THE REASONING LOOP



WORKED EXAMPLE — CONTEXT IN ACTION

“Same weekly grocery order, but half the quantity.”

READ 01
Recognise the pattern
Matches the customer's own permitted order history.

READ 02
Confirm the items
Reads the previous order from the orders service.

READ 03
Apply the change
Quantities halved and validated against stock.

READ 04
Show the new cart
Fresh prices and totals, ready for confirmation.

READ 05
Ask before saving
Any assumption is surfaced as a question first.

How long context is kept

Conversation in session	For the session
Account context	Fetches, not stored
Order and cart data	Read on request
Action record	Retained for audit

Principles we hold to

- ✓ Context is **requested, scoped and perishable** — never a silent profile.
- ✓ Facts come from services; the model contributes language, not data.
- ✓ Ambiguity is resolved by asking, never by guessing on the user's behalf.



Context is fetched, not assumed

Values are read from the platform at the moment of the request, never from memory.



The user can see why

When the agent makes a choice, it can explain which request and which data led to it.



Refusal is a valid outcome

If context or permission is missing, saying so is the correct answer — not a workaround.

“Context is what makes an agent useful. Permission is what makes it safe. InTrust's agent would hold both — and never confuse one for the other.”



4
CONTEXT
SOURCES




6
REASONING
STEPS



0
GUESSED
FACTS

safe by architecture

The model proposes.
The platform decides.



Guardrails, in plain terms


PROPOSED

- | | |
|--------------------------------------|---------------------------------------|
| ✓ No arbitrary SQL | ✓ No unrestricted DB access |
| ✓ No service-role keys in the agent | ✓ No arbitrary API execution |
| ✓ No privilege escalation path | ✓ No trust in model-supplied IDs |
| ✓ Identity resolved server-side | ✓ Explicit capability registry |
| ✓ Permission check before every call | ✓ Business logic stays authoritative |
| ✓ No financial execution capability | ✓ Rate limiting & loop protection |
| ✓ Structured, validated tool results | ✓ Prompt-injection defences |
| ✓ Full audit trail per action | ✓ Graceful, explicit failure handling |

A capability that is not approved simply **does not exist** as far as the agent is concerned.

4 RISK & ACTION MODEL

clear rules per action type



Governance rhythm we would propose

A capability register reviewed with InTrust before anything is enabled, plus a periodic review of agent activity and any user or merchant feedback that suggests a rule needs refining.

Observability we would propose

Every request records the intent, the capability selected, the permission outcome and the result — so behaviour can be examined without changing the platform underneath.

“ An agent is only as safe as the boundary it cannot cross. Here, that boundary is the platform's own permission layer — unchanged, and still in charge.

16
GUARDRAILS

4
RISK TIERS



0
FINANCIAL
ACCESS

ONE LAYER, TWO EXPERIENCES

same brain, different keys

One agent. Two very different users.

Customers and merchants need completely different answers from the same intelligence. What they share is the **foundation**: one identity model, one permission layer, one capability registry, one audit trail. Only the scope of what each is allowed to reach changes.

CUSTOMER EXPERIENCE

shop in your own words

The customer agent

Authenticated as

Customer account

Can reach

Their own orders, cart, wishlist

Typical ask

“Find this, add it, where is my order?”

Writes allowed

Cart & wishlist only

Payment

Not accessible

DISCOVERY

CART

WISHLIST

ORDERS

REWARDS

MERCHANT EXPERIENCE

run your business by asking

The merchant assistant

Authenticated as

Merchant account

Can reach

Their own business data only

Typical ask

“Orders today? Low stock? Earnings?”

Writes allowed

Read-first by default

Panel permissions

Still authoritative

ORDERS

INVENTORY

METRICS

SUBSCRIPTION

SERVICEABILITY

SHARED FOUNDATION · BUILT ONCE

Authentication

Existing sign-in and sessions.

Permissions

Ownership and role rules.

Context layer

Scoped, per-request context.

Capability registry

One reviewed action list.

Guardrails

The same limits apply to both.

Audit & observability

Every action is traceable.

same rules for both

SIDE BY SIDE

One code path · two permission scopes

DIMENSION	CUSTOMER AGENT	MERCHANT ASSISTANT	WHAT IS IDENTICAL
Identity	Customer session	Merchant session	Resolved server-side, never by the model
Data scope	Their own account	Their own business	Scoped by existing permissions
Capabilities	Commerce actions	Operations & insights	Same registry, different subset
Write behaviour	Cart & wishlist	Read-first	Approved before enabling
Payment	Outside the agent	Outside the agent	Never a capability

WHERE THE BOUNDARY SITS

Identical rule for both audiences

Read

Answers and lookups over the user's own data. Safe to enable first.

Write

Cart, wishlist and profile changes — reversible, and revertible.

Never

Payment, credentials and anything final stay outside the agent.

“Two products do not need two intelligences. They need one foundation with the discipline to know exactly how far each user's reach extends.”

2

EXPERIENCES

1

FOUNDATION

1

PERMISSION LAYER



DIGITAL SOLUTIONS FOR A STRONGER TOMORROW


One agent.
Every user.

PHASED ROADMAP

start small, prove it, expand

Phased by design. Every phase stands alone.

Nothing here needs to be delivered at once. Each phase ends with something **usable and observable**, and each one is independently valuable to InTrust. If priorities change, the work to date remains in place.



Sequencing principles

- ✓ **Read before write.** Understanding the data comes first; state changes come later.
- ✓ **Reversible before irreversible.** Cart and wishlist precede anything final.
- ✓ **One audience at a time.** Customer experience is proven before merchant operations.
- ✓ **Capability by capability.** Each addition is reviewed and can be withdrawn.

How we would know a phase worked

- ✓ Requests are being understood and completed without manual intervention.
- ✓ Every action is traceable back to a capability and a permission outcome.
- ✓ Users and merchants complete the intended task more directly than before.
- ✓ Support questions about "how do I..." reduce for the journeys in scope.

What we would need from InTrust

- ✓ A named owner for capability approval — additions signed off by the business, not by the build.
- ✓ Existing role definitions and permission outcomes, so the agent inherits them exactly.
- ✓ Access to the catalogue, stock and order services the panel already uses.
- ✓ A contact for incident and rollback decisions once a phase is live.

Held back on purpose

- ✓ No payment, wallet, OTP or credential handling inside the agent.
- ✓ No pricing, discount or eligibility logic invented by the model.
- ✓ No direct database access — every read and write uses an existing service.
- ✓ No autonomous change to a merchant's records without confirmation.

A design-partner cohort
A small group observed closely in each phase

→ **A feedback loop**
Observations feed the next capability review

→ **A rollback path**
Any capability off, without touching the platform

“ A phase that cannot be measured is a phase that cannot be trusted. Each stage ends with evidence, a reference list of capabilities, and a clear decision to continue. ”

5
PHASES

1
PHASE AT A
TIME


0
BIG-BANG
RELEASES

EXPECTED EXPERIENCE & FUTURE VISION

the long view

Where this leads — an InTrust that participates.

If Phase 01 and Phase 02 land well, the same foundation supports progressively richer behaviour — without a single new permission model, and without the platform ever losing control of what it **is** and what it **is not** allowed to do.



3 A MOMENT IN THE FUTURE STATE

three people, one layer



A customer, on the move

Sends a photo of a bag and says “like this, under ₹3,000”. Gets four matches, adds one to the wishlist and has a reviewed cart waiting — never touching a payment screen in the conversation.



A merchant, before opening the panel

Asks how yesterday went and which items need restocking. Gets both answers from their own records, in seconds, without learning a single new report.



An operations team, quietly

Sees which capabilities are used, which questions recur and where the agent asks instead of acting — using the audit trail rather than guesswork.



What does not change, at any horizon

- ✓ Server-side identity
- ✓ Platform owns the truth
- ✓ Every action auditable
- ✓ Registered capabilities only
- ✓ Payment stays outside
- ✓ Any capability can be withdrawn



How we would read the signals

- ✓ Where requests are clarified too often, the capability or the copy is wrong.
- ✓ Where users still leave the conversation, the hand-off needs designing better.
- ✓ Where the agent refuses, the registry may simply need one more approved action.



What we would measure at each horizon

- ✓ Where requests resolve in one turn, and where they need clarifying.
- ✓ Whether users reach checkout or a decision without leaving the conversation.
- ✓ How often the agent refuses — and whether a missing capability caused it.
- ✓ Whether every action maps back to an approved capability and permission.



What we would deliberately avoid

- ✓ Adding autonomy before the data and permissions beneath it are trusted.
- ✓ Letting the assistant become a second source of business rules.
- ✓ Broadening scope for its own sake — capability only where it earns its place.
- ✓ Any horizon that cannot be switched off cleanly.



Directional vision. Behaviour beyond the pilot phases is illustrative and would be designed, reviewed and approved with InTrust before any build. Nothing described here is committed until it is written into a phase.

“The destination is not a cleverer chatbot. It is an InTrust that quietly completes what its users came to do — and can prove, every time, that it stayed in bounds.



4
HORIZONS



6
FIXED RULES



0
PAYMENT
ACCESS

CLOSING

Talk to InTrust. Get Things Done.

let's build the layer that completes.

InTrust already holds everything this needs: the users, the catalogue, the orders, the merchants, the permissions and the trust. This proposal adds the one thing a conversation still lacks — the ability to **finish what the user started**, inside every control that already exists.

THE PROPOSAL IN ONE LIST

nothing hidden

- ✓ **One intelligence layer** over the platform you already run.
- ✓ **Two experiences** — customer commerce and merchant operations.
- ✓ **A fixed capability registry** — nothing else can be executed.
- ✓ **Existing permissions in charge** — the model never decides access.
- ✓ **Zero payment access** — checkout stays exactly where it is.
- ✓ **Five phases**, each independently valuable and reversible.

5 PROPOSED NEXT STEPS

Sequential · each one confirms the next

Step 01
DISCOVERY

Align on scope, journeys and baselines

Agree the customer and merchant journeys in scope, the events that define completion today, and the numbers we will measure against.

WORKING SESSION

Step 02
DESIGN

Design the architecture and the guardrails

Define the capability registry for Phase 01, the permission mapping, the confirmation rules and the audit requirements.

DESIGN SIGN-OFF

Step 03
BUILD

Build the Phase 01 foundation

The agentic layer, read-only capabilities, permission integration, guardrails and observability — with nothing yet exposed to users.

INTERNAL VALIDATION

Step 04
PILOT

Pilot the customer agent on a contained capability set

Enable discovery, wishlist, cart and order lookup for a controlled group, with the audit trail reviewed throughout.

MEASURED PILOT

Step 05
EXPAND

Measure, refine and widen the registry

Extend to multimodal input and the merchant assistant, adding capabilities only where the evidence supports them.

EVIDENCE-LED GROWTH



One added layer

No rebuild, no parallel backend.



Two experiences

Customers and merchants, one foundation.



Zero payment access

Money moves only in secure checkout.



Governed by default

Permissions, guardrails and audit first.

THE IMMEDIATE NEXT STEP

A working session to agree the Phase 01 capability list.

One conversation is enough to convert this proposal into a scoped, buildable first phase — and to confirm exactly which actions InTrust is comfortable making conversational first.

the InTrust agent



InTrust
A smarter everyday

“ Every journey in this document already exists. The only thing being proposed is a faster, more human way for people to reach them — safely, and on InTrust's terms.



Agentic
CAPABILITY



2
EXPERIENCES



5
PHASES